

Clinical Decision Making and Pattern Recognition in Health Care

Agentic AI, Anomaly Detection, and the Future of Payment Integrity

Submitted by: Rohini Mallikarjunaiah | Stevens Institute of Technology | M.S. Business Intelligence & Analytics

1. Concept Definition

Clinical decision making and pattern recognition in health care encompasses the application of advanced computational methods to support, automate, and augment decisions across Treatment, Payment, and Operations (TPO). At its core, this domain applies machine learning classification, anomaly detection, time-series analysis, and increasingly, agentic generative AI to identify patterns that signal clinical risk, payment fraud, or operational inefficiency.

Agentic AI represents the next evolution beyond passive analytics: systems that autonomously perceive data, reason through evidence chains, and act flagging a suspicious claim, recommending a treatment pathway, or routing an operational exception without requiring a human decision at every step. When combined with supervised classifiers (XGBoost, Random Forest) and unsupervised anomaly detection (Isolation Forest, Autoencoders), these systems create a layered intelligence architecture capable of processing millions of healthcare events in real time.

Cotiviti operates at the intersection of all three TPO domains, serving as the analytics layer between payers and providers. The convergence of agentic AI with clinical pattern recognition directly determines the company's capacity to deliver payment accuracy, reduce waste, and improve health outcomes for its clients.

2. Relevant Trends

Rise of Agentic AI in Healthcare

Large language models (LLMs) have evolved from passive text generators to active reasoning agents capable of executing multi-step clinical workflows. Systems such as Med-PaLM 2 and GPT-4 with tool use demonstrate that LLM agents can interpret clinical notes, cross-reference billing codes, and generate audit-ready narratives tasks previously requiring specialized human reviewers. The global healthcare AI market is projected to reach \$187.95 billion by 2030 [1], driven largely by agentic workflow automation.

Anomaly Detection for Payment Integrity

Healthcare fraud, waste, and abuse cost the U.S. system an estimated \$300 billion annually [2]. Traditional rule-based detection systems catch known fraud patterns but fail against novel schemes. Unsupervised anomaly detection (Isolation Forest, DBSCAN) and gradient-boosted classifiers (XGBoost, LightGBM) now achieve precision rates above 90% on claims fraud detection tasks, significantly outperforming rule engines while reducing false positive rates that burden legitimate providers.

Chain Reasoning and Explainability

Regulatory bodies including CMS and OIG increasingly require explainable AI outputs for audit purposes. Chain-of-thought reasoning in LLMs enables systems to produce human-readable justifications alongside algorithmic flags critical for provider dispute resolution and regulatory compliance. Explainability is no longer optional; it is a market requirement.

Real-Time TPO Integration

The shift from retrospective to prospective payment integrity catching errors before payment rather than clawing back after requires real-time inference pipelines. Edge ML deployment and streaming architectures (Apache Kafka, Flink) now enable sub-second anomaly scoring on live claims feeds, creating new opportunities for pre-payment intervention at scale.

3. Opportunities and Threats

Opportunities

- Pre-payment anomaly interception at scale reduces clawback costs and improves payer-provider relationships
- Agentic AI can automate SAR-equivalent narratives for clinical audits, reducing reviewer workload by 60-80%
- LLM agents trained on payer contracts can autonomously identify billing code mismatches against contracted rates
- Time-series models detecting provider behavior drift can surface emerging fraud schemes before they reach payment thresholds
- Explainable AI outputs create defensible audit trails that satisfy CMS and OIG requirements

Threats

- Adversarial fraud actors increasingly use AI to craft claims that evade ML detection an arms race dynamic
- HIPAA and state privacy regulations constrain the data sharing necessary to train high-quality models
- High false positive rates damage provider relationships and create regulatory exposure under the No Surprises Act
- LLM hallucination in clinical contexts carries patient safety and compliance risk if not rigorously validated
- Consolidation among health tech vendors (Optum, Change Healthcare) creates competitive pressure on independent analytics providers

4. Strategic Recommendations for Cotiviti

Recommendation 1: Deploy an Agentic Claims Review Layer

Cotiviti should invest in an agentic AI layer that sits between ML anomaly detection and human reviewer queues. Rather than surfacing raw flags, the agent performs chain-of-thought reasoning cross-referencing procedure codes against diagnosis codes, contracted rates, provider billing history, and CMS guidelines and delivers a plain-language audit narrative alongside a recommended action (Auto-Approve / Flag for Review / Escalate). This reduces reviewer cognitive load, accelerates decision velocity, and creates a defensible audit trail for each claim disposition.

Recommendation 2: Build a Real-Time Behavioral Anomaly Engine

Cotiviti should move from batch anomaly scoring to real-time provider behavioral profiling using streaming ML pipelines. By tracking rolling 30-day patterns units billed per procedure, duplicate submission rates, multi-provider clustering, submission timing distributions the system can detect emerging fraud signatures before individual claims cross payment thresholds. This shifts Cotiviti's value proposition from retrospective recovery to prospective prevention, a significant competitive differentiator.

Recommendation 3: Develop an Explainability-First AI Governance Framework

As CMS and OIG scrutiny of AI-driven claim denials intensifies, Cotiviti should proactively build an explainability framework that documents model inputs, reasoning chains, and confidence intervals for every AI-assisted decision. This positions Cotiviti as a trusted AI partner rather than a black-box vendor critical for long-term payer client retention and regulatory resilience.

Word Count: ~750 words (excluding bibliography) | Report prepared using primary research, industry publications, and original proof-of-concept development.

Bibliography

- Grand View Research. (2023). "Healthcare Artificial Intelligence Market Size, Share & Trends Analysis Report." Grand View Research. <https://www.grandviewresearch.com/industry-analysis/artificial-intelligence-ai-healthcare-market>
- [2] National Health Care Anti-Fraud Association. (2023). "Health Care Fraud: A Serious and Costly Problem." NHCAA. <https://www.nhcaa.org/tools-insights/about-health-care-fraud/>
- [3] Singhal, K., et al. (2023). "Large Language Models Encode Clinical Knowledge." *Nature*, 620, 172–180. <https://doi.org/10.1038/s41586-023-06291-2>
- [4] Centers for Medicare & Medicaid Services. (2024). "Fraud Prevention System." CMS.gov. <https://www.cms.gov/About-CMS/Components/CPI/Fraud-Prevention-System>
- [5] Liu, X., et al. (2022). "Isolation Forest for Anomaly Detection in Healthcare Claims." *Journal of Biomedical Informatics*, 128, 104-112.
- [6] Cotiviti. (2024). "Payment Integrity Solutions." Cotiviti.com. <https://www.cotiviti.com/solutions/payment-integrity>
- [7] Chen, T., & Guestrin, C. (2016). "XGBoost: A Scalable Tree Boosting System." *Proceedings of KDD 2016*. <https://doi.org/10.1145/2939672.2939785>
- [8] Office of Inspector General. (2023). "Use of Artificial Intelligence in Health Care Fraud Detection." OIG.HHS.gov.